


**DTU**



**Danmarks Tekniske Universitet**

---

# **Colorizing Grayscale Images using UNet**

---

**Lab report  
by**

**Kristian Friis Nielsen (s204120)**

**Mads Andersen (s204137)**

**Iván Viemoes Cuevas (s205823)**

**Anders Nørskov (s183995)**

# Contents

|          |                                                |          |
|----------|------------------------------------------------|----------|
| <b>1</b> | <b>Abstract</b>                                | <b>1</b> |
| <b>2</b> | <b>Introduction</b>                            | <b>1</b> |
| <b>3</b> | <b>Methods and data</b>                        | <b>1</b> |
| 3.1      | Data . . . . .                                 | 1        |
| 3.2      | Architecture . . . . .                         | 1        |
| 3.3      | Reconstruction- and adversarial loss . . . . . | 2        |
| <b>4</b> | <b>Results</b>                                 | <b>3</b> |
| <b>5</b> | <b>Discussion</b>                              | <b>3</b> |
| 5.1      | MNIST results . . . . .                        | 3        |
| 5.2      | Regularization . . . . .                       | 4        |
| 5.3      | Adding depth . . . . .                         | 4        |
| 5.4      | Evaluation and FID . . . . .                   | 4        |
| 5.5      | Datasets . . . . .                             | 5        |
| 5.6      | Generative model . . . . .                     | 5        |
| 5.7      | Adversarial and reconstruction loss . . . . .  | 5        |
| 5.8      | Ethics . . . . .                               | 5        |
| <b>6</b> | <b>Learning Outcome</b>                        | <b>6</b> |
| <b>7</b> | <b>References</b>                              | <b>6</b> |

| Student | 1 | 2 | 3.1 | 3.2 | 3.3 | 4 | 5.1 | 5.2 | 5.3 | 5.4 | 5.5 | 5.6 | 5.7 | 5.8 |
|---------|---|---|-----|-----|-----|---|-----|-----|-----|-----|-----|-----|-----|-----|
| s204120 | X | X |     | X   |     |   | X   | X   | (X) |     |     |     |     |     |
| s204137 | X | X |     |     |     | X |     |     | X   | X   | (X) |     |     |     |
| s205823 | X | X | X   |     |     |   | (X) |     |     |     | X   |     |     | X   |
| s183995 | X | X | (X) | (X) | X   |   |     |     |     |     |     | X   | X   |     |

X denotes primary responsibility. (X) denotes secondary responsibility. However, everyone has contributed to all sections.

## 1. ABSTRACT

Restoring grayscale images, such as historically significant black and white photographs, have been a very manual and time consuming operation. Attempts at automating this process has become a highly contested field in machine learning with only specialized models being used for each unique task. In order to propose a general model we examine the scalability of adversarial networks and their ability to colorize datasets of increasing complexity. We implement a UNet with Instance Normalization as in [1] with adversarial Wasserstein loss [2],  $R_1$  Regularization [3] and Instance Noise [4], and introduce a novel approach called Instance Smoothing much akin to Label Smoothing [5] to stabilize training further. We confirmed the agility of the model across different complexities of datasets, showing that the underlying architecture retained stability and performance as long as the UNet was scaled appropriately, achieving a FID of 16.804 on the *Athletic Fields* category of the Places 365 dataset. The potentials of adversarial networks greatly outweigh their complexity and deliver promising results when trained for extended periods of time with somewhat saturated colorization and without the same brownish tinge affecting results using the simpler MSE loss.

## 2. INTRODUCTION

RGB images are 3D volumes containing color intensities over three channels which are linearly transformed into a grayscale image in a lossy way – that is grayscale images do not contain any hue information, only luminosity, effectively making perfect recoloring a seemingly impossible task. We want to implement a neural network that can reverse this transformation.

Image colorization is a very model intensive task as it must both detect image features and accordingly distribute appropriate color. Additionally, object features can be arbitrarily colored according to a context and pertaining to the object itself which further increases complexity. Different methods for grayscale-colorization using neural networks have been proposed with varying results [6][7]. The objective is not to colorize objects exactly like the ground truth but rather be able to fool the human perception system.

In this report we implement a UNet with instance normalization [1] and build upon this with adversarial networks as in pix2pix [7] without the stochastic element. We implement various methods of stabilizing adversarial network training such as  $R_1$  regularization and instance noise, and introduce a novel approach we call instance smoothing much akin to label smoothing [5]. To achieve perceptually distinguishable images for a human we train the model to predict and discriminate in the CIELAB colorspace.

We hypothesize the adversarial networks to be able to fool a human observer, and to greatly improve the results obtained by conventional colorization models which often use a per-pixel Euclidean distance measure that often leads to a desaturated, brownish tinge.

## 3. METHODS AND DATA

### 3.1. Data

To test our model we use the *MNIST* hand written digit dataset [8] and create a toy dataset to apply to our problem. We colorize the *MNIST* dataset by assigning an interval on the unit circle for each digit and uniformly sample a value from this interval and project it onto a unit square and minmax-normalize it to give us the  $*a$  and  $*b$  channels matching the CIELAB ( $L^*a^*b$ ) colorspace with the definitions in Open Computer Vision 2 [9] for Python with each value ranging from 0 to 1. As such we sample the colors depending on the digit as follows:

$$\begin{aligned} v_\theta &= 2\pi(D + u)/10 \quad \text{with } u \sim \mathcal{U}[-0.5, 0.5] \\ (a_\theta, b_\theta) &= (\sin(v_\theta), \cos(v_\theta)) \\ (*a, *b) &= \left( \frac{0.5a_\theta}{\max(|a_\theta|, |b_\theta|)} + 0.5, \frac{0.5b_\theta}{\max(|a_\theta|, |b_\theta|)} + 0.5 \right) \end{aligned}$$

with  $D \in \{0, 1, 2, \dots, 9\}$  for each digit. The  $L^*a^*b$  color space mimics human perception and spreads out each color such that different values of  $*a$  and  $*b$  are perceptually different. Due to the small range on the unit circle and the sequential uniform distribution of the  $*a$  and  $*b$  channels two sequential digits can be perceptually identical as seen on figure 1.a. The digits vary in quality and some are crudely written and therefore difficult for any model to predict.

In order to evaluate on more complex data, we use the *Athletic Fields* category from the Places 365 dataset [10]. This set contains 40,000 images compromised of both the standard and challenge subsets using the small 256x256 resolution image size for both. We chose the *Athletic Fields* category because it contains easily distinguishable features with mostly uniformly spread well-defined saturated colors focusing only on a single category of places. The *Athletic Fields* dataset most prominently features green terrain but it also regularly varies between red and blue as seen on some images in figure 1.b depending on what sports event is being depicted. Furthermore, it features arbitrarily colored garments and varying skin colors depending on nationality, and what sports team is playing, etc. Many of these arbitrarily colored objects are hard to distinguish in the gray colorspace, since there is some overlap between colors, e.g. yellow, cyan and white etc. due to the linear transformation. All together, this poses a substantially increased challenge for the model without being too complex compared to the whole *Places 365* dataset.

For evaluating results we use the Fréchet inception distance (FID)[11]. This metric uses the Inception v3 [12] model to define a distance between the data distributions of generated and real images - lower distances meaning the generated images have perceptually similar features to the real images.

### 3.2. Architecture

We use the UNet with instance normalization described by [1] and similar to [7] composed of an encoder and decoder network with skip connections from the encoder to the decoder and a bottleneck as shown in figure 2. We use learned



(a.) MNIST Handwritten Digits, colorized



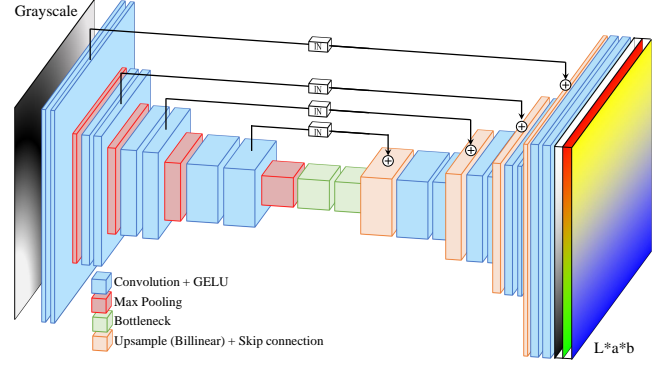
(b.) Places365 Athletic Fields category

**Fig. 1:** The first 64 samples from both datasets.

affine instance normalization on the skip connections. The encoder network detects features while the instance normalized skip connections use the learned *style* to define the colorization of such features in the decoder network. The UNet architecture enables feature detection across scales as the encoder downsamples the input to convolve on a coarser scale while the decoder receives these features and deconstructs them into meaningful colorizations at multiple scales. This architecture fits our problem well since the spatial locations of features in the encoder and decoder are shared.

When training the model we convert each image to grayscale and input it into the network as shown in figure 2. The *MNIST* grayscale images in training are just the originally imported grayscale images. For the sake of consistency with similar literature we call the UNet “the generator  $G(y)$ ” even though the model is fully deterministic according to some input unlike conventional generative networks – that is our model does not depend on a random latent vector  $z$  and as such we write  $G(y)$  instead of  $G(z|y)$  as in [13] for conditional generative networks. For all adversarial networks we use an untrained ResNet [14] model as the discriminator matching the generator in number of parameters.

The image before conversion is then considered the



**Fig. 2:** UNet with instance normalization. The UNet model has 2 channels (\*a,\*b) in its output and the input grayscale image is concatenated as the L channel.

ground truth on which we measure the loss for training.

### 3.3. Reconstruction- and adversarial loss

We introduce a smoother transition for the generator by interpolating between real and generated data when evaluated by the discriminator which we call instance smoothing for image conditional GANs. For each image in a batch we sample  $\alpha_s \sim \mathcal{U}[0, \beta_s]$  decreasing  $\beta_s$  from 0.5 to 0 over a number of iterations. We define two interpolates between real image  $x$  and a generated image  $G(y)$ :

$$\hat{x}_F = \alpha_s x + (1 - \alpha_s) G(y) \quad (1)$$

$$\hat{x}_R = (1 - \alpha_s) x + \alpha_s G(y) \quad (2)$$

$$\text{where } x \sim p_{\text{data}}(x)$$

where  $y$  is the grayscale transformation of  $x$  when training. The loss function then becomes:

$$V(D, G) = \mathbb{E}[f(D(\hat{x}_F|y)) + f(-D(\hat{x}_R|y))]$$

where

$$\begin{aligned} \lim_{\beta_s \rightarrow 0} \mathbb{E}[f(D(\hat{x}_F|y)) + f(-D(\hat{x}_R|y))] \\ = \mathbb{E}[f(D(G(y)|y)) + f(-D(x|y))] \end{aligned}$$

since  $\beta_s$  is the upper bound in the uniform distribution for  $\alpha_s$ . This yields the original conditional GAN minimax game [13] as  $\beta_s$  decreases to zero. This smooth transition in training is motivated by the improvements seen in label smoothing [5] which smoothes the conditional labels of the generator where our method creates a smooth boundary between real and generated data increasing support between the two distributions. For our problem this equates to adding the correct color to the generated image to some degree. During training this addition is then tuned down at which point the generator has hopefully learned to use the correct coloring.

To stabilize training even further we implement instance noise which is defined by adding some noise signal with scale  $\sigma_{\text{critical}}^2 = |f'(0)|/|f''(0)|$  to the input in the discriminator. This further increases support between the real and generated distributions as described in [4].

The gradient penalty for Wasserstein GANs [2] greatly complements our method as the gradient penalty is added to the loss function along interpolated lines between real and generated data:

$$L_{\nabla \hat{x}_F} = \lambda_{\text{recon}} \mathbb{E}_{\hat{x}_F} [(\|\nabla_{\hat{x}_F} D(\hat{x}_F|y)\|_2 - 1)^2]$$

where  $\hat{x}_F$  is the interpolate corresponding to the notation as defined in eq.(1). This is a soft enforcement of the 1-Lipschitz constraint required by the WGAN. This term is similar to our method of instance smoothing as they use the same interpolates. This penalty enforces the gradient to have a maximum norm of one along the interpolated lines. In our case where  $\alpha_s \sim \mathcal{U}[0, \beta_s]$  for decreasing  $\beta_s$  the aforementioned gradient penalty tends to penalize only gradients on generated data and as such we introduce another gradient penalty operating only on real data called  $R_1$  regularization given by:

$$R_{1\nabla x} = \frac{\gamma_{R1}}{2} E_{x \sim p_{\text{data}}(x)} [\|\nabla D(x)\|_2^2]$$

as described in [3] which also is a soft enforcement of the 1-Lipschitz constraint. Using instance noise [4] together with  $R_{1\nabla x}$  regularization is further supported by [15] in which they build upon [2] by introducing exactly this form of the gradient penalty to ensure convergence towards a Nash equilibrium.

As such the final loss function (or minimax game) becomes:

$$\max_G \min_D V(D, G) = \mathbb{E}[f(D(\hat{x}_F|y)) + f(-D(\hat{x}_R|y))] + L_{\nabla \hat{x}_F} + R_{1\nabla x}$$

with  $\hat{x}_R$  and  $\hat{x}_F$  as defined in equation (1) and (2).

We add this adversarial loss to a standard reconstruction loss. In this report we consider Huber (Smooth L1) and Mean Squared Error (L2/MSE) loss. The reconstruction MSE loss is given by:

$$\text{MSE}(x_F, x_R) = \frac{1}{B} \sum_{n=1}^B (x_R - x_F)^2$$

with  $B$  being the batch size. Huber is quadratic like MSE around the origin but linear elsewhere. Huber loss has been shown to improve image saturation [16] and decrease blurring [7].

## 4. RESULTS

We maintained the following parameters throughout the conducted experiments and list different configurations for the two datasets in table 1. The batch size was set to 32, and a learning rate of  $5 \cdot 10^{-5}$  for training on the *MNIST* data set, and  $10^{-4}$  on the *Athletic Fields* data set. For upsampling layers we used *Billinear* and for downsampling layers we used *Maxpool*, this applies to all of our models and on both datasets. We set the  $f(\cdot)$  in the loss function to the softplus function  $f(x) = \log(1 + e^x)$ .

Using the category *Athletic Fields* from the places data set, we randomly crop a 128x128 section of the original 256x256 image each time the image is loaded into a batch for training. For the *MNIST* data we used two decoder and encoder blocks. With the *Athletic Fields* data we used seven encoder and decoder blocks while for configuration M we only used two as in *MNIST*. Configurations K and S were trained without adversarial loss. Configuration P used dropout of 50% applied to the three first decoder layers as described in [7] which makes our model architecture only differ from

*pix2pix* by applying instance normalization on the skip connections. For *MNIST* and *Athletic Fields* we used *ResNet18* and *ResNet152* respectively as the discriminator. All networks were normally initialized with  $\mu = 0$ ,  $\sigma = 0.02$  as described in the DCGAN paper [17]. We calculated the FIDs between 10,000 real and 10,000 generated images to obtain statistically sound results using *pytorch-fid* [18]. All models were trained on a Nvidia Tesla V100 SXM2 32 GB courtesy of DTU Compute.

| Configuration               | GP | Recon. type | $\lambda_{\text{recon}}$ | $\gamma_{R1}$ | InS. <sup>†</sup> | InN. <sup>‡</sup> | FID           |
|-----------------------------|----|-------------|--------------------------|---------------|-------------------|-------------------|---------------|
| A MNIST                     | ✓  | MSE         | 1                        |               |                   | 125               | 15.784        |
| B MNIST                     | ✓  | MSE         | 1                        |               |                   | 1000              | 18.010        |
| C MNIST                     | ✓  | MSE         | 1                        |               | ✓                 | 1000              | 11.790        |
| D MNIST                     | ✓  | MSE         | 100                      |               |                   | 500               | 11.001        |
| E MNIST                     | ✓  | Huber       | 100                      |               | ✓                 | 500               | <b>10.559</b> |
| F MNIST                     | ✓  | Huber       | 1000                     |               |                   | 500               | 10.635        |
| G MNIST                     | ✓  | Huber       | 100                      | 10            | ✓                 | 500               | 10.573        |
| H MNIST                     |    | Huber       | 100                      | 10            | ✓                 | 500               | <b>10.228</b> |
| I MNIST                     |    | Huber       | 100                      | 1             | ✓                 | 500               | 10.550        |
| J MNIST                     |    | Huber       | 100                      | 10            |                   | 500               | 10.301        |
| K MNIST                     |    | Huber       | 100                      |               |                   |                   | 17.613        |
| L Ath. Fields               |    | Huber       | 100                      | 10            | ✓                 | 500               | 25.579        |
| M Ath. Fields - shallow     |    | Huber       | 100                      | 10            | ✓                 | 500               | 38.454        |
| N Ath. Fields               |    | MSE         | 100                      | 10            | ✓                 | 500               | 25.788        |
| O Ath. Fields               |    | Huber       | 100                      | 10            |                   | 500               | 26.142        |
| P Ath. Fields               | ✓  | Huber       | 100                      |               | ✓                 | 500               | 27.127        |
| Q Ath. Fields               | ✓  | Huber       | 100                      | 10            | ✓                 | 500               | <b>25.283</b> |
| R Ath. Fields - pix2pix     |    | Huber       | 100                      | 10            | ✓                 | 500               | 26.252        |
| S Ath. Fields - only recon. |    | Huber       |                          |               |                   |                   | 25.296        |
| X Ath. Fields - 4000 iter.  | ✓  | Huber       | 100                      | 10            | ✓                 | 5000              | 16.804        |

**Table 1:** Hyperparameters and resulting FID scores for a given model trained over 1000 iterations. The models were trained on the same samples, and evaluated on the exact same images. Config. K was trained without the adversarial loss. All *Athletic Fields* models were trained using *MNIST* configuration H as a baseline. Config. M is a shallow architecture with only two blocks in the encoder and decoder. Blank field means configuration without this method. †: Instance Smoothing enabled. ‡: Linearly decrease instance noise to zero over this number of iterations.

Note that all models with adversarial loss in table 1 use instance noise. When training the adversarial networks without instance noise the models often crashed in an early iteration.

## 5. DISCUSSION

### 5.1. MNIST results

The *MNIST* dataset contains both neat- and crudely written digits. During the learning phase the model has to both detect features and colorize them appropriately. While the overwhelming majority of generated *MNIST* images are very close to the ground truth, poorly executed digits without common features are unevenly colored such as col. 10, 16 and 18 in figure 3 as the model is unsure about the exact digit-defining features. The most troublesome digits are those that are similarly indecipherable for a human, e.g. col. 19 and 20. Tricky digits such as col. 2, 10, 13 and 14 show configuration A’s tendency to produce more varied results although being closer to the ground truth in col. 18.


We ascertain that the model learns to colorize only certain digit-defining features individually and struggles to make a unanimous colorization on the digit as a whole. As such some digits are unevenly colored as evident by col. 17 in figure 3 where the upper arc matches a three or a two and the lower part is only present in sevens.

Overall, configurations E, G, H, J produce visually superior results which correlates well with their FIDs. Due to the





**Fig. 3:** Random digits from 0 to 9 and some handpicked digits for the given model configurations.



**Fig. 4:** Random samples from the *Athletic Fields* set and the generated image for the given model configurations on the left. Model L follows the best configuration from MNIST, M is a shallow version of L, model Q is the best performing model on *Athletic fields*, model S is trained without adversarial loss, model X is equivalent to configuration Q but trained over 4000 iterations

visual similarities between configurations K (without adversarial loss) and A (baseline) in figure 3 an unoptimized model is highly dependent on fine tuning before being able to produce the expected improvements from adversarial loss. This is further corroborated by configuration B from table 1 which shows that a wrongly tuned model underperforms compared to the baseline configuration A.

## 5.2. Regularization

Assessing table 1 adding instance smoothing and noise decay generally results in a significant FID decrease. This confirms the positive effect of instance smoothing as the generator learns to utilize more varying color choices matching real images and thus challenging the discriminator. Instance smoothing appears to significantly decrease FID as evident in configurations C and L compared to B and O respectively. However, comparing L and P  $R_1$  regularization decreases FID similarly with instance smoothing. This implies that as long as either instance smoothing or  $R_1$  regularization is used the performance and stability should improve significantly.

As evident in table 1 we achieve the best results implementing both gradient penalties on the *Athletic Fields* dataset. Both gradient penalties are soft enforcements of the 1-Lipschitz constraint. However, they differ in which data distribution to enforce the gradient penalty upon. By using both we hope to more consistently constrain the gradients of the discriminator and stabilize training. An argument can be made that the penalties do not give high enough payback from the added computational complexity. In the worst case they might also counteract each other close to equilibrium.

## 5.3. Adding depth

When ramping up to the *Athletic Fields* dataset the fully convolved UNet clearly yields the best results. This is evident when comparing the shallow model M, which uses the same amount of decoder and encoder blocks as the MNIST models, to its fully featured counterparts (e.g. model L) which otherwise share the exact same parameters. This is most likely due to the fact that we detect coarser features when increasing network depth which is allowed by the increased image resolution. Figure 4 acts as a direct visual representation of this phenomenon. Directly comparing models L (baseline) and M (shallow) the deeper network show more saturated colors particularly in large patches such as the sky or ground. Coloring is also more well defined on items such as clothing and skin. The shallow network in comparison is greatly challenged by the trees breaking up the sky, indicating a poorer ability to separate image features, as seen in col. 5. The disadvantage of increased computation time and model complexity is greatly outweighed by the superior, less bland and detailed results.

In this regard col. 6 in figure 4 works as a good perceptual measurement of the performance of the model. This image has complicated features and a deceiving perspective but still being easily conceivable by the human eye – which boils down to the model’s ability to color inside the lines of the running track which model X and Q clearly does the best. Model X also performs well on arbitrarily colored objects and makes a believable colorization of the jersey in col. 1 even though the ground truth jersey is a solid blue.

## 5.4. Evaluation and FID

When evaluating the results of colorization some image features are more context sensitive than others. Our objective

is to colorize an image in such a way that it does not remove context from the image. Objects which might have arbitrary colors *independent* of their context (e.g. T-shirts, cars) should be given a color that makes sense for the object itself – while objects that have colors that are *dependent* on context (e.g. grass, asphalt, person, sky) should be colored in a way that makes sense in a given context. Higher precision on context-dependent features with solid coloring is preferred, which are often features taking up big proportions of the image thereby increasing significance for human evaluation. Accurate colorization of smaller objects is less significant than for larger more persistent objects.

Making direct pixel-wise measurements on the performance of the model fully ignores context and can easily lead to problematic results. To this extent we utilized the FID metric which is consistent with human judgment as explored in [11] which found that human evaluation is often biased. This metric not only considers the correct color value but also takes into account the individual features of the image. The FID evaluation is more efficient than human judgement and can evaluate more images giving a better estimate of how well the model generalizes.

All models but X trained over 1000 iterations with a batch size of 32 which is 32,000 images out of the 40,000 images in the dataset meaning the models did not run for even a single epoch. No adversarial networks were plateauing or converging yet. This implies that there is room for improvement on all models without overfitting, which is also evident by model X. However, due to time limits it was not feasible to train all the models for any longer.

## 5.5. Datasets

As seen in figure 4 model X col. 2 our model predicts the asphalt as being a red running track. We suppose this is due to the detection of a runner, which in most cases in *Athletic Fields* are on a red surfaces, creating the impression that all runners are on red running tracks. In addition, asphalt is also underrepresented in the dataset which indicates that the trained model will not generalize well outside of *Athletic Fields* settings. By introducing greater variation in the dataset the model would have to train more to achieve the same degree of specificity and still be able to generalize.

The chosen datasets (*MNIST* and *Athletic fields*) were selected because of the difference in complexity. Further research of the scalability of the model on even more complex datasets could be explored. This should be centered around configuration H and Q which performed the best on *MNIST* and *Athletic Fields* respectively. We conclude that the performance of using gradient penalty is dependent on the complexity of the dataset.

## 5.6. Generative model

In the development phase we tried implementing a way to incorporate Gaussian noise in the model to make it truly generative. However, the output ended up being highly deterministic and the generator learned to just ignore the noise – which is consistent with [7] and [19] for image conditional GANs. As a consequence we ended up removing all stochastic input for all models except model R which we trained according to the *pix2pix* model [7] using dropout of  $p = 0.5$  on the first

half of the decoder blocks. This model underperformed its predecessor model L which means the stochastic element of dropout only had a negative effect on the model.

For further research we propose using a mapping network from a latent vector  $z$  to an intermediate latent space  $\mathcal{W}$  which controls the generator by supplying the adaptive instance normalization layers in the decoder with styles as used in StyleGAN [20]. In their paper the affine parameters of the instance normalization layers are stochastically mapped (adaptive) instead of being a learned constant as in our model. Learned constant affine parameters essentially lock the colorization to certain styles.

Using this proposed approach would make the stylization stochastic and generative compared to our constant learned stylization parameters – this would in turn allow the model to predict the color of items which can be arbitrarily colorized since the style input of the instance normalization layer is now also stochastic in nature. However, this would conflict with using a reconstruction loss such as MSE or Huber as those would penalize the model for choosing an arbitrary color that does not match a given image even though it could have been colored a different way without removing context from the image. We propose using a reconstruction loss that is independent of the hue meaning only the correct color saturation is considered although this might lead to problems with the model predicting that grass is blue which leaves it up to the discriminator to discriminate the difference.

## 5.7. Adversarial and reconstruction loss

In early development we found that the model usually converged to a brownish, desaturated solid color for the whole image when training a simple model using only MSE loss – which is consistent with [6], which is why we consider adversarial networks at all. In figure 4 the results for model X are somewhat saturated but still brownish in color in col. 2 and 3. However, it generally colorizes grass and skies with the correct saturation according to ground truth.

Model S without adversarial loss performed marginally worse than the best adversarial model Q according to FID metric in table 1. This could mean that the reconstruction loss has not yet plateaued on the more complex dataset *Athletic Fields*. We expect the reconstruction loss to plateau at some point while the adversarial loss would have more potential to improve the model. This is further corroborated by the results on the simpler dataset *MNIST* where the adversarial networks clearly outperformed the non-adversarial one. Since the models were still not converging after about one epoch, it is hard to predict if or when this will happen since the reconstruction loss would still be present when it plateaus. The adversarial loss would have to outweigh the reconstruction loss in order for the generator to learn.

## 5.8. Ethics

Since our objective is not to be exact, colorizing greyscale images such as historically significant photos could lead to loss of important specific details. Failing to colorize with historical accuracy increases risk of wrongful revisionism. On the other hand it could also contribute to an increase in public interest and immersion into historical context.

## 6. LEARNING OUTCOME

We have learned about the agility of convolutional neural networks and how much customization you can do to them in order to tailor them to your specific needs. We also found out that there are an overwhelming amount of possible parameters to tweak, from the structure itself to the hyperparameters. We have gained a much more robust understanding of the inner workings of convolution layers and how they detect image features - as well as the interesting implications of the skip-connection driven UNet architecture for image recognition tasks. Furthermore, we exposed ourselves to the vast potential of GANs and the adversarial interplay between generators and discriminators as a more advanced means of implementing network loss - for overcoming the shortcomings of MSE. We gathered experience on how images are represented as a 3D datastructure, and how they can be manipulated in python. We also got acquainted with how images can be defined in a colorspace.

## 7. REFERENCES

- [1] Zheng Xu, Xitong Yang, Xue Li, and Xiaoshuai Sun, "The effectiveness of instance normalization: a strong baseline for single image dehazing," *CoRR*, vol. abs/1805.03305, 2018.
- [2] Ishaan Gulrajani, Faruk Ahmed, Martín Arjovsky, Vincent Dumoulin, and Aaron C. Courville, "Improved training of wasserstein gans," *CoRR*, vol. abs/1704.00028, 2017.
- [3] Lars Mescheder, Sebastian Nowozin, and Andreas Geiger, "Which training methods for gans do actually converge?," in *International Conference on Machine Learning (ICML)*, 2018.
- [4] Casper Kaae Sønderby, Jose Caballero, Lucas Theis, Wenzhe Shi, and Ferenc Huszár, "Amortised MAP inference for image super-resolution," *CoRR*, vol. abs/1610.04490, 2016.
- [5] Tim Salimans, Ian J. Goodfellow, Wojciech Zaremba, Vicki Cheung, Alec Radford, and Xi Chen, "Improved techniques for training gans," *CoRR*, vol. abs/1606.03498, 2016.
- [6] Richard Zhang, Phillip Isola, and Alexei A. Efros, "Colorful image colorization," *CoRR*, vol. abs/1603.08511, 2016.
- [7] Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, and Alexei A. Efros, "Image-to-image translation with conditional adversarial networks," *CoRR*, vol. abs/1611.07004, 2016.
- [8] Yann LeCun and Corinna Cortes, "MNIST handwritten digit database," 2010.
- [9] Itseez, "Open source computer vision library," <https://github.com/itseez/opencv>, 2015.
- [10] Bolei Zhou, Agata Lapedriza, Aditya Khosla, Aude Oliva, and Antonio Torralba, "Places: A 10 million image database for scene recognition," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2017.
- [11] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, Günter Klambauer, and Sepp Hochreiter, "Gans trained by a two time-scale update rule converge to a nash equilibrium," *CoRR*, vol. abs/1706.08500, 2017.
- [12] Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jonathon Shlens, and Zbigniew Wojna, "Rethinking the inception architecture for computer vision," 2015.
- [13] Mehdi Mirza and Simon Osindero, "Conditional generative adversarial nets," *CoRR*, vol. abs/1411.1784, 2014.
- [14] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun, "Deep residual learning for image recognition," 2015.
- [15] Naveen Kodali, Jacob D. Abernethy, James Hays, and Zsolt Kira, "How to train your DRAGAN," *CoRR*, vol. abs/1705.07215, 2017.
- [16] Yi Xiao, Peiyao Zhou, and Yan Zheng, "Interactive deep colorization with simultaneous global and local inputs," *CoRR*, vol. abs/1801.09083, 2018.
- [17] Alec Radford, Luke Metz, and Soumith Chintala, "Unsupervised representation learning with deep convolutional generative adversarial networks," 2016.
- [18] Maximilian Seitzer, "pytorch-fid: FID Score for PyTorch," <https://github.com/mseitzer/pytorch-fid>, August 2020, Version 0.1.1.
- [19] Michael Mathieu, Camille Couprie, and Yann LeCun, "Deep multi-scale video prediction beyond mean square error," 2016.
- [20] Tero Karras, Samuli Laine, and Timo Aila, "A style-based generator architecture for generative adversarial networks," *CoRR*, vol. abs/1812.04948, 2018.



# Code: Training on MNIST and Athletic Fields

Also found on: [https://github.com/andersxa/ColorizingGrayscaleImages/blob/main/Colorizing\\_Grayscale\\_Images\\_Colored.ipynb](https://github.com/andersxa/ColorizingGrayscaleImages/blob/main/Colorizing_Grayscale_Images_Colored.ipynb)

```
1 from IPython.display import display, clear_output
2 from matplotlib import pyplot as plt
3 #plt.style.use('dark_background')
4 import numpy as np
5 import torch
6 import cv2
7 from tqdm import tqdm
8 import torch.nn as nn
9 import torch.nn.functional as F
10 import torchvision
11 import torchvision.transforms as transforms
12 import torchvision.models as models
13 from glob import glob
14 import os
15 #from PIL import Image
16 #from skimage import color
17 torch.manual_seed(0)
18 np.random.seed(0)
19
20 # %%
21 #@title Parameters {display-mode: "form"}
22 #@markdown ---
23 batch_size = 64 #@param {type: "number"}
24 disc_iters = 5 #@param {type: "number"}
25 gen_iters = 1 #@param {type: "number"}
26 plot_iter = 50 #@param {type: "number"}
27 #@markdown ---
28 D_lr = 5e-5 #@param {type: "number"} #1e-4 er optimalt
29 G_lr = 5e-5 #@param {type: "number"} #1e-4 er optimalt
30
31 beta_1 = 0.9 #@param {type: "number"}
32 beta_2 = 0.999 #@param {type: "number"}
33 G_betas = D_betas = (beta_1, beta_2) #ONLY FOR ADAM
34 #@markdown ---
35 training_opt = "GP" #@param ["None", "R1_reg", "GP", "R1_GP"]
36 R1_reg = 0.0 #@param {type: "slider", min: 0.0, max: 100.0, step: 0.1}
37 instance_noise_iter = 1000 #@param {type: "slider", min: 0.0, max: 20000, step: 100}
38 instance_smoothing = False #@param {type: "boolean"}
39 #@markdown ---
40 #Training options for reconstruction loss
41 recon_type = 'mse' #@param ["None", "huber", "mse"]
42 recon_lambda = 1.0 #@param [0.1, 0.2, 0.2, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 1.1, 1.2, 1.5, 2.0, 5.0, 10.0, 25.0, 50.0, 100.0, 200.0] {type: "slider", min: 0.1, max: 200.0, step: 0.1}
43 #@markdown ---
44 downsample_type = "MaxPool" #@param ["Conv", "MaxPool"]
45
46 model_name = "MNIST-1mse-no-smooth-1000ins-noise"
47
48 # %%
49 hyper_params = {}
50 hyper_params['batch_size'] = batch_size
51 hyper_params['disc_iters'] = disc_iters
52 hyper_params['gen_iters'] = gen_iters
53 hyper_params['plot_iter'] = plot_iter
54 hyper_params['D_lr'] = D_lr
55 hyper_params['G_lr'] = G_lr
56 hyper_params['beta_1'] = beta_1
57 hyper_params['beta_2'] = beta_2
58 hyper_params['training_opt'] = training_opt
59 hyper_params['R1_reg'] = R1_reg
60 hyper_params['instance_noise_iter'] = instance_noise_iter
61 hyper_params['instance_smoothing'] = instance_smoothing
62 if recon_type in ['mse', 'huber']:
63     hyper_params['recon_type'] = recon_type
64     hyper_params['recon_lambda'] = recon_lambda
65 hyper_params['downsample_type'] = downsample_type
66
67 # %%
68 # For athletic field
69 # min_size = 256
70 # crop_size = 128
71 # if not os.path.exists(f'./{min_size}_sized_images.txt'):
72 #     with open(f'./{min_size}_sized_images.txt', 'w') as f:
73 #         for p in tqdm(glob('data_256/a/athletic_field/outdoor/*.jpg')):
74 #             print(p, file=f)
75
76 # import multiprocessing
77 # resize_crop_transform = transforms.Compose([transforms.Resize(min_size), transforms.RandomCrop(crop_size)])
78 # class CustomDataset(torch.utils.data.Dataset):
79 #     def __init__(self, transform):
80 #         with open(f'./{min_size}_sized_images.txt', 'r') as f:
81 #             self.image_paths = f.read().splitlines()
82 #             self.transform = transform
83
84 #     def __len__(self):
85 #         return len(self.image_paths)
86
87 #     def __getitem__(self, idx):
```

```

88 #         im = cv2.imread(self.image_paths[idx])
89 #         #im_RGB = torch.tensor(cv2.cvtColor(im, cv2.COLOR_BGR2RGB) / 255.0, dtype=torch.float).permute(2,0,1)
90 #         im_Lab = torch.tensor(cv2.cvtColor(im, cv2.COLOR_BGR2Lab) / 255.0, dtype=torch.float).permute(2,0,1)
91 #         #im = self.transform(torch.cat([im_Lab, im_RGB],0))
92 #         im = self.transform(im_Lab)
93 #         return im
94
95 # dataset = CustomDataset(transform=resize_crop_transform)
96 # data_loader = torch.utils.data.DataLoader(dataset,
97 # batch_size=batch_size, shuffle=True, num_workers=multiprocessing.cpu_count(), pin_memory=True)
98 # print(len(dataset))
99
100 # %%
101 import multiprocessing
102
103
104 class ColorMNIST(torch.utils.data.Dataset):
105
106     def __init__(self, dataset):
107         self.data = dataset.data
108         self.targets = dataset.targets
109
110     def __len__(self):
111         return len(self.data)
112
113     def __getitem__(self, idx):
114         pi = 3.1415927410125732
115         image, label = self.data[idx], self.targets[idx]
116         image = image / 255.0
117         label = float(label) + torch.empty(1).uniform_(-0.5, 0.5)
118         a_const = torch.sin(2 * pi * label / 10.0)
119         b_const = torch.cos(2 * pi * label / 10.0)
120         sample = torch.stack([
121             image,
122             0.5 * a_const / max(abs(a_const), abs(b_const)) * image + 0.5,
123             0.5 * b_const / max(abs(a_const), abs(b_const)) * image + 0.5
124         ])
125         return sample
126
127
128 dataset = ColorMNIST(
129     torchvision.datasets.MNIST('data', train=True, download=True))
130 data_loader = torch.utils.data.DataLoader(dataset,
131 batch_size=batch_size,
132 shuffle=True,
133 num_workers=0,
134 pin_memory=True)
135 print(len(dataset))
136
137
138 # %%
139 def get_samples(batch_size):
140     idxs = np.random.randint(len(dataset), size=(batch_size,))
141     return idxs, torch.stack([dataset[i] for i in idxs], 0)
142
143
144 # %%
145 def Lab2RGB(im):
146     if im.shape[0] == 3:
147         im = im.permute(1, 2, 0)
148         im = 255.0 * im
149         im = im.to(torch.uint8).numpy()
150         return cv2.cvtColor(im, cv2.COLOR_Lab2RGB)
151
152
153 # %%
154 plt.figure(figsize=(4.25, 4.25), dpi=180)
155 plt.imshow(Lab2RGB(
156     torchvision.utils.make_grid([dataset[i] for i in range(8 * 8)],
157                                 8).cpu()),
158     interpolation='nearest')
159 plt.axis('off')
160
161
162 # %%
163 def weights_init(m):
164     classname = m.__class__.__name__
165     if classname.find('Conv2d') != -1:
166         nn.init.normal_(m.weight.data, 0.0, 0.02)
167
168
169 class ConvBlock(torch.nn.Module):
170
171     def __init__(self,
172 in_channels,
173 out_channels,
174 n_intermediate_layers=1,
175 direction='in'):
176         super(ConvBlock, self).__init__()
177         if direction == 'out':
178             self.BlockLayers = nn.Sequential(
179                 nn.Conv2d(in_channels,
180 in_channels,

```

```

181         kernel_size=3,
182         stride=1,
183         padding=1), nn.GELU(), *[
184             1 for _ in range(n_intermediate_layers)
185             for l in [
186                 nn.Conv2d(in_channels,
187                         out_channels,
188                         kernel_size=3,
189                         stride=1,
190                         padding=1),
191                 nn.GELU()
192             ]
193         ])
194     else:
195         self.BlockLayers = nn.Sequential(
196             nn.Conv2d(in_channels,
197                     out_channels,
198                     kernel_size=3,
199                     stride=1,
200                     padding=1), nn.GELU(), *[
201                 1 for _ in range(n_intermediate_layers)
202                 for l in [
203                     nn.Conv2d(out_channels,
204                             out_channels,
205                             kernel_size=3,
206                             stride=1,
207                             padding=1),
208                     nn.GELU()
209                 ]
210             ])
211
212     def forward(self, x):
213         x = self.BlockLayers(x)
214         return x
215
216
217 class EncoderBlock(torch.nn.Module):
218
219     def __init__(self,
220                 in_channels,
221                 out_channels,
222                 n_intermediate_layers=1):
223         super(EncoderBlock, self).__init__()
224         self.EncoderBlock = ConvBlock(in_channels,
225                                       out_channels,
226                                       n_intermediate_layers,
227                                       direction='in')
228
229         self.DownSample = nn.MaxPool2d(
230             2) if downsample_type == 'MaxPool' else nn.Sequential(
231                 nn.Conv2d(out_channels,
232                         out_channels,
233                         kernel_size=3,
234                         stride=2,
235                         padding=1), nn.GELU())
236
237     def forward(self, x):
238         x = self.EncoderBlock(x)
239         skip = x
240         x = self.DownSample(x)
241         return x, skip
242
243 class DecoderBlock(torch.nn.Module):
244
245     def __init__(self,
246                 in_channels,
247                 out_channels,
248                 n_intermediate_layers=1):
249         super(DecoderBlock, self).__init__()
250         if upsample_type == 'PixelShuffle':
251             self.UpSample = PixelShuffleUpsample(in_channels, 2) #
252         else:
253             self.UpSample = nn.Upsample(scale_factor=2,
254                                         mode='bilinear',
255                                         align_corners=False)
256         self.DecoderBlock = ConvBlock(in_channels,
257                                       out_channels,
258                                       n_intermediate_layers,
259                                       direction='out')
260
261     def forward(self, x, skip):
262         x = self.UpSample(x)
263         x = self.DecoderBlock(x + skip)
264         return x
265
266
267 # %%
268 #Unet Basic
269 class GeneratorUNet(torch.nn.Module):
270
271     def __init__(self):
272         super(GeneratorUNet, self).__init__()
273         #Encoder Mapping

```

```

274 self.enc_block1 = EncoderBlock(1, 256, 1)
275 self.enc_block2 = EncoderBlock(256, 256, 1)
276
277 #Skip connection Instance Norms
278 self.skip_IN1 = nn.InstanceNorm2d(256, affine=True)
279 self.skip_IN2 = nn.InstanceNorm2d(256, affine=True)
280
281 #Bottleneck
282 self.bot_block1 = ConvBlock(256, 512, 1, direction='in')
283 self.bot_block2 = ConvBlock(512, 256, 1, direction='out')
284
285 #Decoder
286 self.dec_block1 = DecoderBlock(256, 256, 1)
287 self.dec_block2 = DecoderBlock(256, 256, 1)
288
289 #Out
290 self.out_conv = nn.Conv2d(256,
291                             2,
292                             kernel_size=1,
293                             stride=1,
294                             padding=0,
295                             bias=True)
296
297 def forward(self, gray):
298     x = gray
299     #x = torch.cat([x, torch.randn(x.size(0), 2, x.size(2), x.size(3), device=x.device)], 1)
300     x, skip1 = self.enc_block1(x)
301     x, skip2 = self.enc_block2(x)
302
303     #Bottleneck
304     x = self.bot_block1(x)
305
306     #Skip AdaINs
307     skip1 = self.skip_IN1(skip1)
308     skip2 = self.skip_IN2(skip2)
309
310     x = self.bot_block2(x)
311
312     #Decoder
313     x = self.dec_block1(x, skip2)
314     x = self.dec_block2(x, skip1)
315
316     #Out
317     out = torch.sigmoid(self.out_conv(x))
318     out = torch.cat([gray, out], 1)
319     return out
320
321
322 GNet = GeneratorUNet().to(device)
323 GNet.apply(weights_init)
324 print(GNet)
325 print("Generator trainable parameters:",
326       sum(p.numel() for p in GNet.parameters() if p.requires_grad))
327
328
329 # %%
330 #Unet Basic
331 class DiscriminatorNet(torch.nn.Module):
332
333     def __init__(self):
334         super(DiscriminatorNet, self).__init__()
335         #resnet = models.resnet34()
336         #self.resnet = nn.Sequential(resnet.conv1, resnet.bn1, resnet.relu, resnet.maxpool,
337         # resnet.layer1, resnet.layer2, resnet.layer3, resnet.layer4, resnet.avgpool)
338
339         # #Encoder
340         self.enc_conv1 = nn.Conv2d(3,
341                                     128,
342                                     kernel_size=3,
343                                     stride=1,
344                                     padding=1)
345         self.enc_conv2 = nn.Conv2d(128,
346                                     128,
347                                     kernel_size=3,
348                                     stride=1,
349                                     padding=1)
350         self.enc_down1 = nn.Conv2d(128,
351                                    256,
352                                    kernel_size=3,
353                                    stride=2,
354                                    padding=1)
355
356         self.enc_conv3 = nn.Conv2d(256,
357                                    256,
358                                    kernel_size=3,
359                                    stride=1,
360                                    padding=1)
361         self.enc_conv4 = nn.Conv2d(256,
362                                    256,
363                                    kernel_size=3,
364                                    stride=1,
365                                    padding=1)
366         self.enc_down2 = nn.Conv2d(256,

```

```

367         512,
368         kernel_size=3,
369         stride=2,
370         padding=1)
371
372     self.enc_conv5 = nn.Conv2d(512,
373                                512,
374                                kernel_size=3,
375                                stride=1,
376                                padding=1)
377     self.enc_conv6 = nn.Conv2d(512,
378                                512,
379                                kernel_size=3,
380                                stride=1,
381                                padding=1)
382
383     #Out
384     self.out_conv = nn.Conv2d(512,
385                                1024,
386                                kernel_size=1,
387                                stride=1,
388                                padding=0)
389
390     def forward(self, x):
391         x = res = F.gelu(self.enc_conv1(x)) #
392         x = res + F.gelu(self.enc_conv2(x)) #
393         x = res = F.gelu(self.enc_down1(x))
394
395         x = res = res + F.gelu(self.enc_conv3(x)) #
396         x = res = res + F.gelu(self.enc_conv4(x)) #
397         x = res = F.gelu(self.enc_down2(x))
398
399         x = res = res + F.gelu(self.enc_conv5(x)) #
400         x = res + F.gelu(self.enc_conv6(x)) #
401
402         x = F.gelu(self.out_conv(x))
403         return x
404
405
406 DNet = DiscriminatorNet().to(device)
407 DNet.apply(weights_init)
408 print(DNet)
409 print("Discriminator trainable parameters:",
410       sum(p.numel() for p in DNet.parameters() if p.requires_grad))
411
412
413 # %%
414 def Lab2RGB(im):
415     if im.shape[0] == 3:
416         im = im.permute(1, 2, 0)
417         im = 255.0 * im
418         im = im.to(torch.uint8).numpy()
419         return cv2.cvtColor(im, cv2.COLOR_Lab2RGB)
420
421
422 def smooth(scalars, weight):
423     last = scalars[0]
424     smoothed = list()
425     for point in scalars:
426         smoothed_val = last * weight + (1 - weight) * point
427         smoothed.append(smoothed_val)
428         last = smoothed_val
429     return smoothed
430
431
432 def anneal(val, target):
433     return max((target - val) / target, 0)
434
435
436 # %%
437 from torch.cuda.amp.autocast_mode import autocast
438 #from torch.cuda.amp import GradScaler
439
440 from collections import defaultdict
441
442 G_optimizer = torch.optim.Adam(GNet.parameters(),
443                                 lr=G_lr,
444                                 betas=G_betas)
445 D_optimizer = torch.optim.Adam(DNet.parameters(),
446                                 lr=D_lr,
447                                 betas=D_betas)
448
449 losses_print = defaultdict(lambda: [])
450 print_to_legend = {
451     'Div.': r'Div.',
452     'G.i+1': r'$D(G_{i+1}(z|y))$',
453     'D.': r'$D(x)$',
454     'G.i': r'$D(G_i(z|y))$'
455 }
456
457 mse_loss_fn = nn.MSELoss()
458 huber_loss_fn = nn.SmoothL1Loss(beta=0.5)
459 iteration = -1

```

```

460 #Linear annealing
461
462 # %%
463
464 with tqdm(range(1000),
465           unit_scale=(gen_iters + disc_iters) * batch_size,
466           unit='img',
467           ncols=200) as tqdm_bar: #, autocalc():#
468     for iteration in tqdm_bar:
469         GNet.train()
470         DNet.train()
471         #https://arxiv.org/pdf/1610.04490.pdf & https://arxiv.org/pdf/1701.04862.pdf
472         instance_noise_annealing = anneal(iteration, instance_noise_iter)
473         DNet.zero_grad()
474         #---start disc loop---
475         disc_real_avg = 0.0
476         disc_fake_avg = 0.0
477         for _ in range(disc_iters):
478             idxs, real = get_samples(batch_size)
479             real = real.to(device).requires_grad_(True)
480             gray = real[:, :1]
481
482             fake = GNet(gray)
483
484             #Training Discriminator
485
486             if instance_smoothing:
487                 alpha = 0.5 * instance_noise_annealing * torch.rand(
488                     real.size(0), 1, 1, 1, device=real.device).expand_as(real)
489                 real_input = (1 - alpha) * real + alpha * fake.detach()
490             else:
491                 real_input = real
492
493             real_noise = (instance_noise_annealing *
494                          np.sqrt(2)) * torch.randn_like(real,
495                                                          device=device)
496
497             disc_real = DNet(real_input + real_noise)
498
499             Dreal_loss = F.softplus(-disc_real).mean()
500
501             if instance_smoothing:
502                 alpha = 0.5 * instance_noise_annealing * torch.rand(
503                     real.size(0), 1, 1, 1, device=real.device).expand_as(real)
504                 fake_input = alpha * real + (1 - alpha) * fake.detach()
505             else:
506                 fake_input = fake.detach()
507
508             fake_noise = (instance_noise_annealing *
509                          np.sqrt(2)) * torch.randn_like(real,
510                                                          device=device)
511
512             disc_fake = DNet(fake_input + fake_noise)
513
514             Dfake_loss = F.softplus(disc_fake).mean()
515
516             #https://arxiv.org/pdf/1801.04406.pdf - R1 Regularization
517             grad_penalty = 0.0
518             if training_opt in ['R1_reg', 'R1_GP']:
519                 real_grad = torch.autograd.grad(outputs=disc_real.sum(),
520                                                  inputs=real,
521                                                  create_graph=True,
522                                                  only_inputs=True)[0]
523                 grad_penalty += (R1_reg / 2) * real_grad.view(
524                     real_grad.size(0), -1).square().sum(-1).mean()
525
526             #https://arxiv.org/pdf/1704.00028.pdf afsnit 4 - gradient penalty
527             if training_opt in ['GP', 'R1_GP']:
528                 alpha = torch.rand(real.size(0), 1, 1, 1,
529                                   device=real.device).expand_as(real)
530                 interp = alpha * real + (1 - alpha) * fake.detach()
531                 disc_interp = DNet(interp)
532                 interp_grad = torch.autograd.grad(outputs=disc_interp.sum(),
533                                                  inputs=interp,
534                                                  create_graph=True,
535                                                  only_inputs=True)[0]
536                 grad_penalty += 10.0 * F.relu(
537                     interp_grad.view(interp_grad.size(0), -1).norm(
538                         2, dim=1).subtract(1.0)).square().sum(-1).mean()
539
540             D_loss = Dreal_loss + Dfake_loss + grad_penalty
541             D_loss.backward()
542
543             disc_real_avg += disc_real.mean().item()
544             disc_fake_avg += disc_fake.mean().item()
545         #---end disc loop---
546         D_optimizer.step()
547         losses_print['D.'].append(disc_real_avg / disc_iters)
548         losses_print['G.i'].append(disc_fake_avg / disc_iters)
549
550         #Training Generator
551         GNet.zero_grad()
552         #---Start generator loop---

```



```

553 disc_gen_avg = 0.0
554 for _ in range(gen_iters):
555     idxs, real = get_samples(batch_size)
556     real = real.to(device).requires_grad_(True)
557     gray = real[:, :1]
558
559     fake = GNet(gray)
560
561     if instance_smoothing:
562         alpha = 0.5 * instance_noise_annealing * torch.rand(
563             real.size(0), 1, 1, 1, device=real.device).expand_as(real)
564         gen_input = alpha * real.detach() + (1 - alpha) * fake
565     else:
566         gen_input = fake
567
568     gen_noise = (instance_noise_annealing *
569                 np.sqrt(2)) * torch.randn_like(real, device=device)
570
571     disc_gen = DNet(gen_input + gen_noise)
572
573     G_loss = F.softplus(-disc_gen).mean(
574         ) #criterion(output, torch.full_like(output, 1.0, dtype=torch.float, device=device))
575
576     recon_loss = 0.0
577     if recon_type == 'mse':
578         recon_loss = mse_loss_fn(fake[:, 1:], real[:, 1:])
579     elif recon_type == 'huber':
580         recon_loss = huber_loss_fn(fake[:, 1:], real[:, 1:])
581     (G_loss + recon_lambda * recon_loss).backward()
582
583     disc_gen_avg += disc_gen.mean().item()
584 #---end generator loop---
585 G_optimizer.step()
586 loss_print['G.i+1'].append(disc_gen_avg / gen_iters)
587
588 with torch.no_grad():
589     tqdm_bar.set_postfix_str(
590         f"Iteration {iteration}, I.N.A.: {instance_noise_annealing:.3f}, "
591         + ", ".join([
592             f"{ns}: {ls[-1]:.4f}" for ns, ls in loss_print.items()
593         ]))
594     if iteration % plot_iter == 0:
595         clear_output(wait=True)
596         #evaluate
597         fig, axs = plt.subplots(4, 11, figsize=(22, 8), dpi=100)
598         gs = axs[0, 5].get_gridspec()
599         for axs_y in axs[:, 5:]:
600             for ax in axs_y:
601                 ax.remove()
602         fig.suptitle(f"{model_name}, Iteration {iteration}, " +
603                     ", ".join([
604                         f"{k}: {v}" +
605                         ("\n" if (i + 1) % 7 == 0 else "")
606                         for i, (k,
607                             v) in enumerate(hyper_params.items())
608                     ]))
609         axbig = fig.add_subplot(gs[:, 5:])
610         axbig.tick_params(direction='in', pad=-22)
611         legend = []
612         for ns, ls in loss_print.items():
613             axbig.plot(ls)
614             legend.append(print_to_legend[ns])
615         axbig.legend(legend)
616         im_pred_ = fake[:, -3:].cpu().detach()
617         im_gray_ = gray.cpu().detach()
618         im_true_ = real[:, -3:].cpu().detach()
619         for i in range(4):
620             axs[i][0].imshow(
621                 Lab2RGB(im_pred_[i].permute(1, 2, 0).float() *
622                     torch.tensor([0.0, 1.0, 1.0]).view(1, 1, 3) +
623                     torch.tensor([0.5, 0.0, 0.0]).view(1, 1, 3)),
624                 interpolation='nearest',
625                 aspect='equal',
626                 vmin=0,
627                 vmax=1)
628             axs[i][1].imshow(
629                 Lab2RGB(im_true_[i].permute(1, 2, 0).float() *
630                     torch.tensor([0.0, 1.0, 1.0]).view(1, 1, 3) +
631                     torch.tensor([0.5, 0.0, 0.0]).view(1, 1, 3)),
632                 interpolation='nearest',
633                 aspect='equal',
634                 vmin=0,
635                 vmax=1)
636             axs[i][2].imshow(im_gray_[i].permute(1, 2,
637                 0).squeeze().float(),
638                             cmap='gray',
639                             interpolation='nearest',
640                             aspect='equal',
641                             vmin=0,
642                             vmax=1)
643             axs[i][3].imshow(Lab2RGB(im_pred_[i].permute(1, 2,
644                 0).float()),
645                             interpolation='nearest',

```

```

646         aspect='equal',
647         vmin=0,
648         vmax=1)
649     axs[i][4].imshow(Lab2RGB(im_true_[i].permute(1, 2,
650                                         0).float()),
651                     interpolation='nearest',
652                     aspect='equal',
653                     vmin=0,
654                     vmax=1)
655     axs[i][0].set_axis_off()
656     axs[i][1].set_axis_off()
657     axs[i][2].set_axis_off()
658     axs[i][3].set_axis_off()
659     axs[i][4].set_axis_off()
660     axs[i][2].set_title(f"ID:{idxs[i]}")
661     #fig.savefig(f"I_{iteration}_results.png", dpi=180)
662     plt.show()
663     #torch.save({'gnet':GNet.state_dict(), 'dnet':DNet.state_dict()}, 'save_athletic_field.h')
664
665 # %%
666 torch.save({
667     'gnet': GNet.state_dict(),
668     'dnet': DNet.state_dict()
669 }, f'{model_name}.h')
670
671 # %%
672 s_dict = torch.load(f"{model_name}.h")
673
674 # %%
675 s_dict = torch.load(f"{model_name}.h")
676 #GNet.load_state_dict(s_dict['gnet'])
677
678 # %%
679 for i in range(20):
680     real_sample = get_samples(512)
681     for i_r, r in enumerate(real_sample):
682         im_rgb = Lab2RGB(r)
683         cv2.imwrite(f"MNIST_real/{i*512+i_r}.jpg", im_rgb)
684
685 # %%
686 get_ipython().run_cell_magic('bash', '-s "$model_name"', 'mkdir "$1"')
687
688 # %%
689 real_paths = list(glob("MNIST_real/*.jpg"))
690 with torch.no_grad():
691     for i in range(20):
692         im_paths = real_paths[i * 512:(i + 1) * 512]
693         gray = torch.stack([
694             torch.tensor(cv2.imread(f, cv2.IMREAD_GRAYSCALE),
695                             device=device).unsqueeze(0) / 255.0
696             for f in im_paths
697         ])
698         fake_sample = GNet(gray)
699         for i_f, f in enumerate(fake_sample):
700             im_rgb = Lab2RGB(f.cpu())
701             cv2.imwrite(f"{model_name}/{os.path.basename(im_paths[i_f])}",
702                         im_rgb)
703
704 # %%
705 get_ipython().run_cell_magic(
706     'bash', '-s "$model_name"',
707     'python3 -m pytorch_fid --device cuda:0 "MNIST_real" "$1"')
708
709 # %%
710 model_name

```